

# Lecture A for foreign students

## Lecture 1

---

Norbert Pozar (npozar@se.kanazawa-u.ac.jp)

Apr 19, 2018

# Today's plan

- Basic `gcc` toolchain usage
- Using the command line efficiently
- Shell scripts for automation

- The GNU Compiler Collection
- Website: <https://gcc.gnu.org/>
- Current version: 7.3
- Supports C (**gcc**), C++ (**g++**), Fortran (**gfortran**), ...
- Alternative: clang (used on macOS by default)

# Basic usage

Compile and link:

```
$ gcc main.c -o main  
$ g++ main.cpp -o main
```

Run the resulting binary:

```
$ ./main
```

**main.c, main.cpp** input source file (**.c** for C and **.cpp** for C++)

**-o main** compiler flag **-o** with a parameter for the name of the produced executable<sup>1</sup>

<sup>1</sup>On Unixy systems (Linux, macOS), executables do not have an extension by convention. On Windows it would be **main.exe**.

# Optimizations

- By default, the compiler **does not perform any optimizations**.
- Important for development and debugging  $\rightsquigarrow$ 
  - Original structure of the code is preserved.
  - It takes time to optimize code.
- Optimization can **significantly speed up** the code: Use optimized builds for running computations!

```
$ gcc main.c -o main -O3
```

**-O3** Turns on the highest level of optimizations<sup>2</sup>.

---

<sup>2</sup>See gcc manual: 3.10 Options That Control Optimization

## Example: Sum $1 + \dots + 10^9$ (optimization1.c)

```
#include <stdio.h>

int main()
{
    double sum = 0.;
    for (int i = 0; i < 1000000000; i++) {
        sum += i;
    }
    return 0;
}
```

Build it and run it (with timing)

```
$ gcc optimization1.c -o optimization1
$ time ./optimization1
```

## Run with optimizations enabled

Build it and run it (with timing)

```
$ gcc optimization1.c -o optimization1 -O3  
$ time ./optimization1
```

# Optimization surprises

- Optimized binary should produce the same *output* as the nonoptimized one.
- Code that does not influence the output can be safely removed: This is one of the most important optimizations!



Let's make sure that we print the result and try again.

```
#include <stdio.h>

int main()
{
    double sum = 0.;
    for (int i = 0; i < 10000000000; i ++) {
        sum += i;
    }

    printf("%f\n", sum);

    return 0;
}
```

## More optimizations: `-march=native`

- Instructs the compiler to use any instructions that your processor supports, including the modern ones.
- A good idea to always use this. (It is unlikely that you will try to run the same binary on another machine without recompiling.)

```
$ gcc optimization2.c -o optimization2 -O3 \
    -march=native
```

## More optimizations: -ffast-math

Numbers on the computer behave differently from real numbers.

**Example.** Addition is not associative. On a computer

$$(a + b) + c \neq a + (b + c)$$

With double precision, try  $a = 10^{16}$ ,  $b = -a$ ,  $c = 1$ .

**How is  $a + b + c$  computed?**

Compiler follows the order of operations and emits code for  $(a + b) + c$ .

## More optimizations: `-ffast-math`

But using  $a + (b + c)$  might be more efficient:

- Maybe  $b + c$  was already computed.
- Modern processors are **superscalar**: Can execute multiple **independent** instructions per cycle.

$$a + b + c + d$$

may be possibly computed faster as

$$(a + b) + (c + d)$$

since  $a + b$  and  $c + d$  can potentially be computed **at the same time**.

→ instruction-level parallelism

## More optimizations: `-ffast-math`

Unfortunately, changing the order of floating point operations can change results.

⇒ Compiler cannot reorder operations by default.

**`-ffast-math`** Flag that tells compiler that it is **OK** to reorder floating point operations<sup>3</sup> to produce a more efficient code, even if it changes the result.

```
$ gcc optimization2.c -o optimization2 -O3 \
    -ffast-math
```

`-O3 -ffast-math` can be replaced by `-Ofast`

---

<sup>3</sup>And a few other optimizations.

## More optimizations: `-ffast-math`

### Use with great care!

Most of the time OK but some numerical algorithm are very sensitive to the order of operations. Test if `-ffast-math` changes your results.

`-ffast-math` also assumes that all floating point results are finite numbers, so no NaNs and no infinities. Make sure your code does not rely on these.

It also cheats a bit on some operations like division, using faster but less accurate methods.

# General optimization tips for faster code

- Start with choosing an **efficient algorithm for your problem**.
  - bubble sort VS quicksort
  - Gaussian elimination VS iterative methods
- **Only ever optimize “hot code”**: code that actually is run the most in the program.
- Use timing (or even profiling<sup>4</sup>) tools to test if changes are actually beneficial and to find which code to focus on.

---

<sup>4</sup>**perf** on Linux: See **perf** Examples for more info. This is quite advanced topic.

## Quick tip for numerical code

**Division** is much slower than multiplication.

If you need to divide by a constant in a hot loop, precompute the reciprocal  $c = 1/b$  outside of the loop and replace  $a/b$  with  $a * c$ .



# Useful Unix commands

**time command args...** Run `command args...` and print the time used.

**top, htop** Show CPU and memory usage of running processes.

**perf** Advanced profiling on Linux. Very detailed information down to which instructions in your code take the most time.

## Enable more warnings

-Wall -Wextra flags enable many useful warnings that help identify serious bugs in your code early. **Always use.**

```
#include <stdio.h>
int main() {
    double x;
    printf("%f\n", x); // uninitialized variable
}

int n(int i) {
    i; // forgotten return + no effect
}
```